

# Case Study

## Data Analysis on Device Dataset

Muhammad Rasyid Ridho Husein


<https://github.com/muhammadrasyidrh/Generated-Device-Dataset>

# Problem Overview



## Objective

Data analysis on device utilization dataset using three different using three different methods



## Tool

Python



## Data


Generated data about device usage data and location type



# Dataset Information

Generated dataset consists of 1000 rows and 10 columns

1. User ID : Integer
2. Location : String
3. Login Time : Date Time
4. Name : String
5. Email : String
6. Phone Number : String
7. Birth Year : Integer
8. Cellphone Brand : String
9. Digital Interest : String
10. Location Type : String



# Dataset Overview

User ID		Location	Login Time		Name \
0	1	Port Johnathan	2024-01-22	22:42:21.083785	Matthew Ballard
1	2	North Bailey	2024-08-31	15:09:29.617381	Nicole Green PhD
2	3	Ronaldberg	2024-07-22	09:22:54.121453	William Ramos
3	4	West Amanda	2024-08-31	13:54:22.955823	Dr. Michael Meyer
4	5	Sarafort	2024-10-13	19:29:12.397854	Micheal Griffith
5	6	New Jamiebury	2024-03-25	22:14:39.023823	Justin Thomas
6	7	Port Henry	2024-11-20	08:36:34.133092	Stephen Willis
7	8	West Danielfurt	2024-09-23	09:48:05.289213	Candice Turner
8	9	South Leahton	2024-06-26	14:21:46.026850	Richard Smith
9	10	Shelleyport	2024-11-19	17:24:05.512022	Steven Cook

	Email	Phone Number	Birth Year	Cellphone Brand \
0	<a href="mailto:matthew.ballard@outlook.com">matthew.ballard@outlook.com</a>	389-3270	2002	Oppo
1	<a href="mailto:nicole.green@yahoo.com">nicole.green@yahoo.com</a>	585-3245	2000	Oppo
2	<a href="mailto:william.ramos@yahoo.com">william.ramos@yahoo.com</a>	323-7142	1999	Samsung
3	<a href="mailto:dr..michael@yahoo.com">dr..michael@yahoo.com</a>	784-5647	1986	Vivo
4	<a href="mailto:micheal.griffith@yahoo.com">micheal.griffith@yahoo.com</a>	565-4150	2003	Huawei
5	<a href="mailto:justin.thomas@yahoo.com">justin.thomas@yahoo.com</a>	109-4301	1980	Xiaomi
6	<a href="mailto:stephen.willis@yahoo.com">stephen.willis@yahoo.com</a>	162-1852	1990	Xiaomi
7	<a href="mailto:candice.turner@yahoo.com">candice.turner@yahoo.com</a>	294-4551	1996	Xiaomi
8	<a href="mailto:richard.smith@yahoo.com">richard.smith@yahoo.com</a>	993-7612	1993	Oppo
9	<a href="mailto:steven.cook@outlook.com">steven.cook@outlook.com</a>	560-1618	2000	Oppo

	Digital Interest	Location Type
0	Travel	Urban
1	Shopping	Suburban
2	Music	Rural
3	Tech News	Urban
4	Sport	Suburban
5	Shopping	Rural
6	Gaming	Urban
7	Gaming	Suburban
8	Travel	Rural
9	Sport	Suburban

# Dataset Overview

**Email** Consist of three domains, that are Gmail, Yahoo, and Outlook

**Cellphone Brand** Consist of six brands, that are Apple, Samsung, Huawei, Oppo, Vivo, and Xiaomi

**Digital Interest** Consist of eight interests, that are Gaming, Fitness, Social Media, Music, Tech News, Shopping, Travel, and sports

**Location Type** Consist of three types, that are Urban, Suburban, and Rural

```
[ ] df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1000 entries, 0 to 999
```

```
Data columns (total 10 columns):
```

#	Column	Non-Null Count	Dtype
0	User ID	1000 non-null	int64
1	Location	1000 non-null	object
2	Login Time	1000 non-null	object
3	Name	1000 non-null	object

4	Email	1000 non-null	object
5	Phone Number	1000 non-null	object
6	Birth Year	1000 non-null	int64
7	Cellphone Brand	1000 non-null	object
8	Digital Interest	1000 non-null	object
9	Location Type	1000 non-null	object

dtypes: int64(2), object(8)  
memory usage: 78.2+ KB

# Development of New Variables

- Separation of **Name Variable** into **First Name** and **Last Name**
- Separation of **Email Variable** into **Email Username** and **Email Domain**
- Separation of **Phone Number Variable** into **Area Code** and **Local Number**
- Separation of **Login Time Variable** into **Login Hour** and **Login Date**
- Established **Age Variable** using **Birth Year Variable**

```
[ ] df['Age'] = 2024 - df['Birth Year']
```

```
[ ] df['Login Time'] = pd.to_datetime(df['Login Time'])
df['Login Hour'] = df['Login Time'].dt.hour
```

```
[ ] # Generate Name
df['First Name'] = df['Name'].apply(lambda x: x.split()[0])
df['Last Name'] = df['Name'].apply(lambda x: x.split()[1])
```

```
[ ] # Generate Email
df['Email Username'] = df['Email'].apply(lambda x: x.split('@')[0])
df['Email Domain'] = df['Email'].apply(lambda x: x.split('@')[1])
```

```
[ ] # Generate Phone Number
df['Area Code'] = df['Phone Number'].apply(lambda x: x.split('-')[0])
df['Local Number'] = df['Phone Number'].apply(lambda x: x.split('-')[1])
```

```
[ ] # Generate Login Time
df['Login Date'] = df['Login Time'].dt.date
```

# Confidence Interval

Applied on variables consist of numeric data.

$$\text{Formula : } CI = \text{mean} \pm Z_{\alpha/2} \times \frac{\text{Std.Error}}{\sqrt{n}}$$

```
[53] # Confidence Interval of Age, Birth Year, and Login Hour
age_lower, age_upper = confidence_interval_mean(df['Age'], confidence=0.95)
birthyear_lower, birthyear_upper = confidence_interval_mean(df['Birth Year'], confidence=0.95)
loginhour_lower, loginhour_upper = confidence_interval_mean(df['Login Hour'], confidence=0.95)

print(f"95% Confidence Interval for Age: ({age_lower:.2f}, {age_upper:.2f})")
print(f"95% Confidence Interval for Birth Year: ({birthyear_lower:.2f}, {birthyear_upper:.2f})")
print(f"95% Confidence Interval for Login Hour: ({loginhour_lower:.2f}, {loginhour_upper:.2f})")
```

```
→ 95% Confidence Interval for Age: (30.95, 31.91)
95% Confidence Interval for Birth Year: (1992.09, 1993.05)
95% Confidence Interval for Login Hour: (11.23, 12.10)
```

# Preprocessing Data

Drop unused variables on the models.

```
[ ] df.drop(['Name', 'User ID', 'Location', 'Login Time', 'Email', 'Phone Number', 'Birth Year', 'First Name', 'Last Name',  
            'Email Username', 'Area Code', 'Local Number', 'Login Date'], axis=1, inplace=True)
```

Remaining variables

```
[ ] df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1000 entries, 0 to 999
```

```
Data columns (total 6 columns):
```

#	Column	Non-Null Count	Dtype
0	Cellphone Brand	1000 non-null	object
1	Digital Interest	1000 non-null	object
2	Location Type	1000 non-null	object
3	Age	1000 non-null	int64
4	Login Hour	1000 non-null	int32
5	Email Domain	1000 non-null	object

```
dtypes: int32(1), int64(1), object(4)
```

```
memory usage: 43.1+ KB
```



# Preprocessing Data

Get dummies using **One Hot Encoder** on **Predictor Variables** and **Ordinal Encoder** on **Response Variable**

**X** = ['Cellphone Brand', 'Digital Interest', 'Age', 'Login Hour', 'Email Domain']

**y** = ['Location Type']

```
[ ] X = pd.get_dummies(df, columns=['Cellphone Brand', 'Digital Interest', 'Email Domain']).drop('Location Type', axis=1)
```

```
[ ] y = df['Location Type']  
    ordinal_encoder = OrdinalEncoder(categories=[['Rural', 'Suburban', 'Urban']])  
    y = ordinal_encoder.fit_transform(y.values.reshape(-1, 1))
```

Split data into training data (80% of data) and testing data (20% of data)

```
[ ] X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```


```
[ ] X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
→ ((800, 19), (200, 19), (800, 1), (200, 1))
```

# Statistical Model



**Logistic  
Regression**



**Random Forest**



**Support Vector  
Machine**



# Logistic Regression



## Training Data

```
[ ] from sklearn.linear_model import LogisticRegression
    from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
    from sklearn.metrics import mean_absolute_error, mean_squared_error, root_mean_squared_error
```

```
[ ] LR = LogisticRegression()
```

```
[ ] LR.fit(X_train, y_train)
```

```
[ ] print('Coef : ', LR.coef_)
    print('Intercept : ', LR.intercept_)
```

```
Coef : [[ 6.90283082e-03  1.00708679e-04  5.03659858e-02 -6.65259390e-02
 -6.48589094e-02 -6.76759075e-02 -4.10579691e-02  7.47721745e-02
 -5.27431643e-02 -1.98425091e-01 -1.38669242e-01  9.88308782e-02
 4.71910475e-02  5.12556276e-02  1.54187914e-02  6.21605883e-02
 7.65870334e-02 -1.03899071e-01 -8.76685272e-02]
[-5.45373805e-03 -1.41819573e-02  5.91021188e-02 -3.45886360e-02
 1.35923504e-01  1.06823206e-01  3.53041422e-02 -9.40991538e-02
 -1.60544619e-01  2.44414978e-01  1.44594265e-03  1.20861821e-01
 -1.32532697e-01  3.07818828e-01 -3.60728147e-02 -1.36926259e-01
 4.41485583e-02  3.57553455e-02  1.28561277e-01]
```

```
[-1.44909277e-03  1.40812486e-02 -1.09468105e-01  1.01114575e-01
 -7.10645943e-02 -3.91472985e-02  5.75382694e-03  1.93269792e-02
 2.13287783e-01 -4.59898868e-02  1.37223299e-01 -2.19692699e-01
 8.53416494e-02 -3.59074456e-01  2.06540233e-02  7.47656703e-02
 -1.20735592e-01  6.81437255e-02 -4.08927499e-02]]
Intercept : [-0.12466047  0.22595597 -0.1012955 ]
```

# Logistic Regression



## Comparison of Training Data and Testing Data

### Training Data

MAE : 0.79875

MSE : 1.19875

RMSE : 1.0948744220229094

### Training Data

Accuracy : 0.40125

### Confusion Matrix :

[[ 88 99 84]

[ 68 129 75]

[ 76 77 104]]

### Classification Report :

	precision	recall	f1-score	support
0.0	0.38	0.32	0.35	271
1.0	0.42	0.47	0.45	272
2.0	0.40	0.40	0.40	257
accuracy			0.40	800
macro avg	0.40	0.40	0.40	800
weighted avg	0.40	0.40	0.40	800

### Testing Data

MAE : 0.855

MSE : 1.255

RMSE : 1.120267825120404

### Testing Data

Accuracy : 0.345

### Confusion Matrix :

[[15 22 19]

[19 30 21]

[21 29 24]]

### Classification Report :

	precision	recall	f1-score	support
0.0	0.27	0.27	0.27	56
1.0	0.37	0.43	0.40	70
2.0	0.38	0.32	0.35	74
accuracy			0.34	200
macro avg	0.34	0.34	0.34	200
weighted avg	0.34	0.34	0.34	200

Accuracy : 0.345

# Random Forest

Random Forest using Hyperparameter Tuning to obtain the best parameters.

```
[ ] # Grid Search CV
    param_grid = {
        'n_estimators': [100, 200, 300, 400, 500],
        'max_depth': [3, 5, 7, 10],
        'min_samples_leaf': [5, 10, 25, 50, 100]
    }

[ ] rf2 = GridSearchCV(estimator=rf1, param_grid=param_grid, cv=5, verbose=1, n_jobs=-1)

[ ] rf2.fit(X_train, y_train)
```

Best Parameters for Random Forest Model on Training Data

- Max Depth of 10
- Min Samples leaf of 5
- Number of estimators of 400
- Accuracy of 0.351249

# Random Forest

## Testing Data

Testing Data

Accuracy : 0.33

Confusion Matrix :

```
[[18 25 13]
 [21 32 17]
 [31 27 16]]
```

Classification Report :

	precision	recall	f1-score	support
0.0	0.26	0.32	0.29	56
1.0	0.38	0.46	0.42	70
2.0	0.35	0.22	0.27	74
accuracy			<u>0.33</u>	200
macro avg	0.33	0.33	0.32	200
weighted avg	0.33	0.33	0.32	200

MAE = 0.89

MSE = 1.33

RMSE = 1.1532562594670797

# Support Vector Machine



## Training Data

Support Vector Machine using Hyperparameter Tuning to obtain the best parameters.

```
[ ] # Grid Search
    param_grid_svc = {
        'C': [100, 10, 1, 0.1, 0.01, 0.001, 0.0001],
        'kernel': ['linear', 'poly', 'rbf', 'sigmoid']
    }

[ ] svc2 = GridSearchCV(estimator=svc1, param_grid=param_grid_svc, cv=5, verbose=1, n_jobs=-1)

[ ] svc2.fit(X_train, y_train)
```

Best Parameters for Support Vector Machine Model on Training Data

- Regularization Parameter of 1
- Kernel of Polynomial
- Accuracy of 0.365

# Support Vector Machine



## Testing Data

Testing Data

Accuracy : 0.345

Confusion Matrix :

```
[[ 6 39 11]
 [11 46 13]
 [10 47 17]]
```

Classification Report :

	precision	recall	f1-score	support
0.0	0.22	0.11	0.14	56
1.0	0.35	0.66	0.46	70
2.0	0.41	0.23	0.30	74
accuracy			<u>0.34</u>	200
macro avg	0.33	0.33	0.30	200
weighted avg	0.34	0.34	0.31	200

MAE = 0.76

MSE = 0.97

RMSE = 0.9848857801796105



# Conclusion

Model	Accuracy	MAE	MSE	RMSE
Logistic Regression	0.345	0.855	1.255	1.12026
Random Forest	0.33	0.89	1.33	1.15325
Support Vector Machine	0.345	0.76	0.97	0.98488

- Overall, the accuracy model is not relatively high, at 34,5%.
- Support Vector Machine model is the optimal model because the model has the lowest error compared to Logistic Regression and Random Forest.
- The Support Vector Machine model can be applied to the generated device data using polynomial kernel function and regularization parameter of 1.

# Thanks!

Data Analysis on Device Dataset

Muhammad Rasyid Ridho Husein  
muhammadrasyidrh@gmail.com  
[linkedin.com/in/muhammadrasyidrh](https://www.linkedin.com/in/muhammadrasyidrh)  
[github.com/muhammadrasyidrh](https://github.com/muhammadrasyidrh)

